

TECH CONTINGENT, IIT MADRAS × DRONAHQ PRESENTS



Tech Contingent



Platform Partner

A 51-HOUR BUILDATHON

# Inter Guild Buildathon

## Building the Autonomous SDR

*One Problem Statement, One Integrated Build*

**51 HOURS • ONE BUILD**

STARTS

Fri, 18 Sep 2026, 9:00 PM

ENDS

Sun, 20 Sep 2026, 11:59 PM

PLATFORM PARTNER

**DronaHQ**

Agentic AI • Low-Code • Vibe Coding

## Welcome

---

This is the problem statement for the 51-hour Buildathon: one integrated build, not separate assignments. The goal is a single autonomous Sales Development Representative (SDR) system that can run real go-to-market campaigns across multiple channels, on top of the DronaHQ agentic platform.

The system has two halves that have to work as one product: a **control plane** that lets a human manager create, configure, launch, pause and monitor campaigns, and an **intelligence layer** of agents that actually research, personalise, contact and follow up with prospects inside those campaigns. Neither half is optional, and neither is judged on its own: what matters is whether the two come together into something that behaves like a real SDR.

## 1 Event at a Glance

---

| Format           | Details                                                                              |
|------------------|--------------------------------------------------------------------------------------|
| Structure        | One combined buildathon, one integrated product                                      |
| Duration         | 51 hours                                                                             |
| Starts           | Friday, 18 September 2026, 9:00 PM                                                   |
| Ends             | Sunday, 20 September 2026, 11:59 PM                                                  |
| Platform partner | DronaHQ (Agentic AI, Vibe Coding, Apps Studio), details on the next page             |
| The build        | An autonomous, multi-channel SDR system: control plane plus agentic out-reach engine |

## Our Platform Partner

### DronaHQ

*Deltecs Infotech Pvt. Ltd. • Founded 2007 • Mumbai & San Francisco*

DronaHQ is the platform this Buildathon is built on. It is an AI-powered low-code platform for building internal tools, AI agents and full operational workflows on top of existing data and systems. Under one roof it brings together a drag-and-drop **app builder** for dashboards, forms and admin panels; an **Agentic AI platform** for designing, deploying and orchestrating autonomous agents (chatbots, voice agents, data agents) with built-in tool-calling, memory, RAG and guardrails; and a **Vibe Coding** mode that turns natural-language prompts into working, editable apps and agents.

### Get Started

|                           |                                                                                                                                           |
|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------|
| Apps Studio, sign up      | <a href="https://www.dronahq.com/signup/">https://www.dronahq.com/signup/</a>                                                             |
| Apps Studio, login        | <a href="https://studio.dronahq.com/">https://studio.dronahq.com/</a>                                                                     |
| Agentic Platform, sign up | <a href="https://www.dronahq.com/agents/signup/">https://www.dronahq.com/agents/signup/</a>                                               |
| Agentic Platform, login   | <a href="https://agents.dronahq.com/">https://agents.dronahq.com/</a>                                                                     |
| Documentation             | <a href="https://docs.dronahq.com/agents/getting-started/introduction/">https://docs.dronahq.com/agents/getting-started/introduction/</a> |

### Follow & Tag DronaHQ

While you post about your build, tag **@DronaHQ**.

|             |                                                                                                                     |
|-------------|---------------------------------------------------------------------------------------------------------------------|
| LinkedIn    | <a href="https://www.linkedin.com/company/deltecs-infotech/">https://www.linkedin.com/company/deltecs-infotech/</a> |
| Instagram   | <a href="https://www.instagram.com/dronahq">https://www.instagram.com/dronahq</a>                                   |
| X (Twitter) | <a href="https://x.com/DronaHQ">https://x.com/DronaHQ</a>                                                           |
| Discord     | <a href="https://discord.com/invite/F2pB4UgEeK">https://discord.com/invite/F2pB4UgEeK</a>                           |

**Mandatory:** DronaHQ must be used meaningfully in your core solution: as the Agentic AI/Vibe Coding engine, the control-plane UI, or both. It should not be a wrapper dashboard bolted onto a system built entirely elsewhere.

## 2 The Problem: Building an Autonomous SDR

Sales development today is manual, repetitive and inconsistent: reps research prospects, personalise outreach, chase replies across LinkedIn, email, SMS and calls, and manually log everything to a CRM. Your job over the next 51 hours is to build a system where AI agents do this work autonomously, while a human manager stays fully in control through a clean control plane.

The one question judges will keep coming back to:

*"How close is this to a real SDR working autonomously across multiple channels?"*

The two sections below describe that one system from two angles: how it is managed, and how it thinks. Build them together, not as separate deliverables.

## 3 Campaign Management & Control Plane

### Multiple Concurrent Campaigns

The system must support several campaigns running at once. A **campaign** is an autonomous GTM program aimed at a specific ICP (Ideal Customer Profile), with its own objectives, agents, prompts, policies and channel configuration: for example *US SaaS CTO outreach*, *India BFSI CIO outreach*, *US Voice AI founders*, and *Existing customer expansion*, all running side by side. Each campaign operates independently while sharing common platform infrastructure, integrations and knowledge where it makes sense to.

### Campaign as a First-Class Object

Every campaign needs its own configuration and lifecycle. At minimum, a campaign should carry:

- ▶ **Identity:** name, description, owner, status, created/last-modified dates
- ▶ **Targeting:** ICP, geography, target roles, company criteria, exclusion criteria, sample/reference profiles
- ▶ **Agent configuration:** which agents are enabled, their responsibilities, tools, policies, decision thresholds, escalation rules
- ▶ **System prompts / AI harness:** a campaign-level system prompt plus agent-specific prompts covering qualification, research, personalisation, outreach, follow-up, response-handling, guardrails and escalation

Prompt/harness versions should be tracked and roll-back-able. A change to one campaign's prompts must never silently change another campaign's behaviour.

### Campaign Lifecycle

Draft → Live → Paused → Completed / Archived

- ▶ **Draft:** being configured, not permitted to send autonomous outreach.
- ▶ **Live:** agents may discover, research, qualify and contact prospects per configuration.
- ▶ **Paused:** all autonomous execution stops immediately; existing prospect and conversation data is retained; no new outreach fires; can be resumed later.
- ▶ **Completed / Archived:** no longer active, but its history, decisions and analytics remain fully available.

### Campaign Control

An authorised user needs a simple way to create, edit, activate, pause, resume, complete, archive and duplicate a campaign, and to view its configuration, activity and performance. A prominent **Pause / Resume** control should sit right on the campaign dashboard, and the current state of every campaign should be obvious at a glance, for example:

| Campaign             | ICP                | Status | Prospects | Outreach | Meetings |
|----------------------|--------------------|--------|-----------|----------|----------|
| US SaaS CTO          | SaaS CTO           | • Live | 1,284     | 426      | 18       |
| India BFSI CIO       | BFSI CIO           | Paused | 642       | 211      | 11       |
| Voice AI Founders    | AI Founders        | • Live | 389       | 142      | 9        |
| Enterprise Expansion | Existing Customers | Draft  | 0         | 0        | 0        |

## Campaign Isolation & Conflict Handling

Each campaign's ICP, prompts, agent configuration, outreach policies, targeting rules, activity state, metrics and decision history should be logically isolated. A prospect may appear in more than one campaign, so the system needs to detect and manage duplicate outreach, two campaigns targeting the same prospect at once, conflicting instructions, and excessive contact frequency, with a clear mechanism to resolve these conflicts.

## Campaign Dashboard

Each campaign gets its own operational dashboard covering:

- ▶ **Status:** Live/Paused/Draft, owner, ICP, active channels, active reps
- ▶ **Prospect funnel:** Discovered → Researched → Qualified → Contacted → Engaged → Meeting → Opportunity
- ▶ **Outreach activity:** LinkedIn actions, emails, calls, messages, follow-ups
- ▶ **Agent activity:** active/completed/failed workflows, pending approvals, escalations
- ▶ **Outcomes:** positive/negative responses, meetings, qualified opportunities, conversion rates

## Global vs. Campaign-Level Configuration

|                           |                                                                                                                                                                                                                      |
|---------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Global<br>(platform-wide) | Authentication, integration credentials, available models & tools, org/security policies, common knowledge bases, global suppression / do-not-contact lists, platform-level guardrails                               |
| Campaign-level            | ICP, geography, target personas, objective, prompts, agent instructions, outreach/follow-up strategy, qualification criteria, campaign-specific knowledge, channel configuration, daily limits, human-approval rules |

## Prompt / Harness Management

Provide a dedicated interface where an authorised user can view the campaign system prompt and per-agent prompts, edit them, save a new version, compare versions, activate a version, roll back, and see who changed what and when. The platform should record *which* prompt/harness version was active whenever an agent takes a meaningful action, so managers can always answer “why did the agent behave this way?” and “which configuration produced this outcome?”

## Levels of Operational Control

- ▶ **Campaign Pause:** stops autonomous execution for one campaign.
- ▶ **Agent Pause:** stops one specific agent while the rest of the campaign continues where appropriate.
- ▶ **Channel Pause:** stops activity on one channel only (e.g. LinkedIn or email).
- ▶ **Global Kill Switch:** stops *all* autonomous external actions across the entire platform.

## Duplication & Experimentation (Stretch)

Consider letting a campaign be duplicated into a variant: for example *US SaaS CTO: Campaign A* duplicated into *US SaaS CTO: Personalisation Variant B* with a different prompt, outreach strategy, qualification criteria or channel mix, so teams can compare response rates, positive-response rates, meeting rates, qualification accuracy, cost per qualified prospect and agent efficiency.

## Representative Assignment

Representatives (human sales reps) should be assignable to one or more campaigns, with control over which reps can execute a campaign, whose identity is used for outreach, daily activity limits, working hours and channel availability. If a rep is offboarded, the system should surface every campaign that's affected and let an admin reassign it.

## Required Demonstration

Your final product must demonstrate **at least three concurrent campaigns**, each with a different ICP, different instructions/prompts, a different target audience, and independent Live/Paused state, dashboards, agent execution and analytics, and you must show that pausing one campaign does *not* stop the others.

### Product design questions worth thinking about

How does a manager understand *all* campaigns at a glance? How are Live vs. Paused campaigns visually differentiated? What happens if a campaign is paused mid-execution? How do configuration changes propagate? Should prompt changes need approval? How are conflicting campaigns on the same prospect handled? How do managers compare performance across campaigns? What should a manager see *before* activating a campaign? How can an operator safely stop autonomous activity in a hurry?

## 4 Agentic Intelligence & RAG

This is the other half of the same system: the intelligence that makes the control plane's campaigns actually run. Your solution should show meaningful use of **Agentic AI, NLP, RAG and AI orchestration** to autonomously execute and optimise GTM outreach inside the campaigns defined in Section 3.

### Agents to Build

- ▶ **ICP Fitment Agent:** evaluates, qualifies or rejects prospects against campaign-specific ICP criteria.
- ▶ **Lead Research & Enrichment Agent:** researches companies/people and builds structured prospect context.
- ▶ **Outreach Strategy Agent:** decides whether, when, where and how a prospect should be contacted.
- ▶ **Personalisation / Email Agent:** uses prospect context and relevant examples to generate contextual outreach.
- ▶ **Conversation Agent:** reads email/SMS/LinkedIn replies and decides the next action.
- ▶ **Voice SDR Agent:** conducts calls, qualifies prospects, handles objections, escalates when needed.
- ▶ **Follow-up Agent:** decides when and how to follow up, or when to stop.

### RAG & Knowledge

Build campaign/product-specific knowledge bases covering product & company information, customer case studies, sales playbooks, ICP definitions, objection handling, high-quality example emails/messages for different scenarios, and voice conversation examples/scripts. Agents should retrieve relevant knowledge *before* making an important decision or generating anything customer-facing, rather than hallucinate from the base model alone.

### Prompting, Guardrails & Evaluation

Explore campaign-specific system prompts and agent harnesses; structured outputs and tool calling; guardrails and hallucination prevention; human-in-the-loop escalation; prompt/version management; golden datasets and AI evaluation; LLM-as-judge or human evaluation; model selection and routing; and fine-tuning, only where it's demonstrably worth it.

### Cost & Performance

Design for production-scale usage: token consumption, model cost, latency, context optimisation, RAG optimisation, caching, using smaller models for simpler tasks, and cost per prospect, per qualified lead, and per conversation.

### Suggested Tools (equivalents welcome)

| Area                      | Options                                               |
|---------------------------|-------------------------------------------------------|
| GTM platform              | DronaHQ Agentic AI + Vibe Coding                      |
| Lead discovery/enrichment | Apollo or equivalent                                  |
| LinkedIn                  | Browser automation or third-party services            |
| Email                     | Gmail / Gmail API                                     |
| SMS & telephony           | Twilio                                                |
| Voice AI                  | Twilio + a voice AI platform, or a custom voice stack |
| Knowledge / RAG           | Vector DB: pgvector / Pinecone / Qdrant / equivalent  |
| Agent tools               | MCP / tool calling                                    |
| Multi-agent comms         | A2A or an agent-orchestration framework               |
| CRM                       | Salesforce / HubSpot / Google Sheets / equivalent     |

The goal is *not* to use every tool in this list. It's a modular, cost-efficient, measurable AI system where specialised agents reason over knowledge, use tools, take real actions, and collaborate, and it should feel like **one SDR working across channels**, not five disconnected bots.

## 5 Engineering Expectations

- ▶ **DronaHQ is required, not sufficient.** Use it meaningfully as your Agentic AI, Vibe Coding or control-plane layer, but the entire codebase should not be built *only* on top of the DronaHQ platform. Judges expect real, engineer-authored code underneath it: custom services, integrations, data models and orchestration logic that you designed and wrote, with DronaHQ as a genuine part of the system rather than the whole of it.
- ▶ **Manage the codebase properly.** One shared repository with meaningful, incremental commits from the whole team; a clear folder structure that separates agents, services, data and frontend; secrets and credentials kept out of version control; a README that explains setup, environment variables and architecture; and no dead, duplicated, or half-finished code paths left in by submission time.
- ▶ **Build for reliability.** Handle malformed model output, failed API calls and empty states without crashing. A smaller feature set that works end to end beats a larger one that only works in the demo path.

## 6 Deliverables

At submission, your team needs to hand in four things:

- ▶ **A report, with everything a judge needs.** A written report covering your problem approach, system architecture (with a diagram), the agents and campaigns you built, how DronaHQ is used in the system, the rest of your tech stack, what's fully working versus partially working versus skipped, and any known limitations or trade-offs you made under the 51-hour clock.
- ▶ **A fully deployed, working website.** A live URL, not a local build, that judges can open and use directly: create and manage campaigns, watch agents work, and see the required three concurrent campaigns behaving independently, exactly as described in Section 3.
- ▶ **A GitHub repository, with a proper README.md.** One public repository holding the full codebase, with a README that includes at least: a project overview, setup and run instructions, required environment variables, the tech stack, an architecture overview, and a clear file/folder structure section explaining what lives where and why.
- ▶ **A presentation deck.** A concise PPT that walks through the problem, your approach, the architecture, what you built, and a live or recorded demo of it working, for the pitch/judging round.

A working deployed link and a clean, well-documented repository are how judges verify everything the report claims. Missing either one makes the rest of your submission hard to trust, however good the build.

## 7 How You'll Be Judged

Judges score across 8 weighted categories. This is the shape of the rubric, enough to help you prioritise your 51 hours; the full detailed checklist stays with the judges.

| Category                       | Points | What it's checking, in short                                                                          |
|--------------------------------|--------|-------------------------------------------------------------------------------------------------------|
| Multi-Channel SDR Intelligence | 25     | Does it intelligently coordinate LinkedIn, email, SMS & voice, like one SDR rather than five bots?    |
| DronaHQ Usage (mandatory)      | 15     | Is DronaHQ meaningfully core to the build, not just a UI wrapper?                                     |
| Context & Personalisation      | 15     | Do agents actually understand the prospect before acting, without hallucinating?                      |
| End-to-End SDR Capability      | 15     | How much of Find → Research → Qualify → Contact → Follow-up → Respond → Book/Escalate actually works? |
| Product & User Experience      | 10     | Could a real sales team pick this up and use it?                                                      |
| Software & Engineering Quality | 10     | Architecture, modularity, reliability, error handling, security.                                      |
| Measurement & Optimisation     | 5      | Can you measure and improve the system's own performance?                                             |
| Innovation / Extra Thinking    | 5      | Ideas beyond the obvious implementation.                                                              |

Quality of decision-making and execution beats sheer number of agents or tools used. A tightly integrated system that clearly demonstrates the control plane and the intelligence layer working together will always outscore a pile of disconnected features.

## 8 Final Notes

- ▶ Build **one** coherent product: the control plane and the intelligence layer should feel like a single system, not a demo stitched together at the last minute.
- ▶ DronaHQ's platform, docs and socials (the Platform Partner section) are your fastest path to a working agentic build in 51 hours: start there, then build outward.
- ▶ When you're ready to demo, be ready to show *at least three* concurrent campaigns behaving independently, as described in Section 3.

Good luck, go build the SDR of the future.